




a-i: mean $\pm$ s.e. over 40 seeds (42-81); FreeSolv GP-base EI n = 39 (one failed fit). Budgets as in Fig. 1 except Matbench-gap at 20.

j: paired difference per seed, then averaged; TL rows 20 seeds (42-61).

TL: TL-base = base transfer-learning surrogate · TL-Ptj = pretrain-then-joint · TL-E2E = end-to-end joint · TL-SPS = soft parameter sharing · TL-MMD = domain adaptation (MMD)
 GP: GP-base = autoregressive multi-fidelity GP (MFGP) · GP-SV = sparse variational MFGP · GP-DKL = deep-kernel GP